Adam Ho

408-203-9917 | adamho129@gmail.com | Portfolio | LinkedIn | GitHub

Summary

2+ YOE in Software QA Automation and Data Analysis, across startup and consumer-tech research operations. Skilled in **triaging CI/CD** failures and customer-reported issues, **analyzing diagnostic logs**, and building **automated regression coverage**.

Skills

QA: Test Automation (Playwright) • Triage • Regression Testing • CI/CD • Bug Reporting • Issue Tracking • Git • API Testing • REST • GraphQL • Postman • Technical Documentation • Root-Cause Analysis • SLA-Based Prioritization
Languages/Tools: Python • TypeScript/JavaScript • SQL • R • Java • C++ • pandas • NumPy • BigQuery • Airflow • Docker
AI/Data: scikit-learn • TensorFlow • A/B Testing & Evaluation Design • Hypothesis Testing • Looker • Tableau • Mixpanel • GA4

Experience

QA Wolf Oct 2025 – Oct 2026
Software QA Engineer *Remote*

- **Triaged and root-caused 100s of test failures daily** from CI/CD runs and user-reported tickets across ∼10 concurrent enterprise-scale SaaS clients, working a **prioritized queue under strict SLAs** for suite investigation times.
- Diagnosed recurring and intermittent failures from **CI logs, API responses, browser console/network output, stack traces, and service logs**, isolating **concurrency and process regressions**, environment differences, and flaky tests.
- Built and maintained **Playwright/TypeScript** automation (reusable page objects, API-based setup/teardown helpers via DevTools network calls): **400+ tests created**, including new tests for user-reported issues, scaling coverage to **2,500+ tests per deploy**; onboarded an AI-powered NLP startup with a **CI/CD** pipeline of 100+ tests, eliminating its need for manual QA.
- Acted as primary liaison between client stakeholders and engineers in scoping sessions to **prioritize test coverage** and reduce **flake** (false-positive) rates; wrote stakeholder-facing reports that gave clients **deploy confidence.**
- **Filed bug reports** with documented repro steps, Loom recordings, and console logs in **Jira** and **Linear**; cross-referenced findings against clients' own internal tickets to catch duplicate and misdiagnosed issues.
- Leveraged **Agentic ML workflows** and custom MCP integrations to accelerate triage, debugging, and test authoring.

Apple (Akkodis) Apr 2025 – Oct 2025
User Research Study Coordinator *Sunnyvale, CA*

- Ran user studies on ∼900 participants to compare design and hardware variations. Designed **study protocols**, participant **eligibility logic**, and subjective measurement scales, aligning research design with **UX Research** and **Product** stakeholder goals.
- Built a pipeline consolidating metrics from individual study trackers into an **automated Tableau dashboard** for weekly business reviews across 8 concurrent user studies.
- Performed **QA testing** and delivered UI mockups for internal tools that offload raw study data to S3; reported bugs and logged feature requests in internal tickets.

Google (Cognizant) Oct 2024 – Apr 2025
Data Specialist *Mountain View, CA*

- Took on analytics scope beyond the moderator role: built **Looker** dashboards on **BigQuery SQL** to track prototype-testing metrics, surfacing **participant retention drivers** to Android XR Product Managers.
- **Automated** retroactive **data-entry validation** and test scheduling through Google Cloud Platform integrations, streamlining downstream **ML** model delivery by a projected 30%.
- Developed and maintained **Python** scripts for ingestion, annotation, and analysis of hardware telemetry used in 600+ internal data collection studies.

Projects

Spotify 'Smart Queue' Simulated A/B Experiment | Python, SQL, BigQuery, Mixpanel, Next.js Oct 2026

- Designed an end-to-end evaluation framework for a simulated feature launch, **defining success metrics**, guardrails, and decision rules in advance. Ran **power analysis** and SRM checks, then evaluated a seeded simulation of 14,420 users (398K rows) using Welch, non-inferiority, and bootstrap CACE methods.
- Built a reproducible workflow from Python through BigQuery SQL to the Mixpanel API, and documented the **tracking plan**, repository, and resulting launch recommendation.

Amazon GPU ETL Pipeline | Airflow, Docker, Python, PostgreSQL Feb 2025

- Built a scheduled **ETL** (Airflow DAG in Docker) that scrapes and transforms GPU listings into Postgres for automated ingestion.

Partner Engineering - EDA & Supervised Model | SQL, R, sklearn, pandas Jan 2024 – Jul 2024

- Cleaned and explored 40M+ real estate records; deployed a **Balanced Random Forest** classifier (3rd place, UCI Expo) and contributed to an **ML** report-writing recommendation system (86% accuracy).

Education

University of California, Irvine Sep 2024
B.S. Data Science *Donald Bren School of Information and Computer Sciences*