Adam Ho

408-203-9917 | adamho129@gmail.com | Portfolio | LinkedIn | GitHub

Summary

2+ YOE in Product Analytics and Software QA Automation, across consumer tech research (Apple, Google) and SaaS startup quality operations. Skilled in **experiment design and analysis**, **building data ingestion and testing pipelines**, and in turning numbers into clear recommendations for product and engineering stakeholders.

Skills

Product Analytics: A/B Testing & Experiment Design • Power Analysis • Funnel & Retention Analysis • KPI & Metric Definition • Event Tracking & Implementation • Root-Cause Analysis • Dashboards & Self-Serve Reporting • Mixpanel • GA4 • PostHog
Technical: SQL • BigQuery • Python (pandas, NumPy, SciPy, statsmodels, scikit-learn, TensorFlow) • Databricks • Next.js • TypeScript/JavaScript • R • Looker • Tableau • Power BI • Excel • Git • REST API • GraphQL • Postman • Airflow • Docker

Experience

QA Wolf Oct 2025 – Oct 2026
Software QA Engineer *Remote*

- **Led** QA automation for an **AI analytics startup**: enabled **product launches** with automated **regression testing** of multiple **new features**, defined **success criteria**, served as main **stakeholder liaison**, and resolved greenlight and deploy issues.
- Built **CI/CD pipelines** to sharply cut flake (false-positive) rates by **80%** and gave **deploy confidence** in major product changes. Scoped **test coverage** and flake-reduction priorities by risk and impact across ∼12 concurrent enterprise SaaS clients.
- Drove **requirements gathering**, translating **acceptance criteria** into structured test plans for product and engineering stakeholders; automated **500+ tests via Playwright/TypeScript**, scaling team coverage to **2,500+ tests per deploy**.
- **Triaged and root-caused 100+ test failures daily** from CI/CD runs and user-reported tickets, working a **prioritized queue under strict SLAs**; documented repro steps, Loom recordings, and console logs in **Jira** and **Linear**, cross-referencing clients' internal tickets to catch duplicate and misdiagnosed issues.
- Developed **custom MCP integrations** connecting our AI-powered QA platform and Playwright runners to **agentic Claude Code workflows**, accelerating triage, debugging, and test authoring.

Apple (Akkodis) Apr 2025 – Oct 2025
User Research Analyst *Sunnyvale, CA*

- Designed and analyzed **A/B experiments** comparing UI and hardware variants across a ∼900-participant VR product study, using **NumPy/pandas** for **non-parametric hypothesis testing** and **SQL** to deliver **data-driven recommendations** to product.
- Defined and reported **KPIs** across 8 concurrent user studies, consolidating metrics from individual trackers into an **automated Tableau dashboard** that became the source of truth for weekly cross-functional business reviews and **prioritization decisions**.
- Designed **experiment protocols**, participant **segmentation and eligibility logic**, and subjective measurement scales, aligning research design with **UX Research** and Product goals; translated stakeholder feature requests into tracked **requirements**.

Google (Cognizant) Oct 2024 – Apr 2025
Data Analyst *Mountain View, CA*

- Translated business questions into **product metrics** for Android XR prototype studies, surfacing **participant retention drivers** to Product Managers via **Looker** dashboards on **BigQuery SQL** that informed prototype decisions.
- Built reusable **SQL segmentation** and eligibility logic on a database of 600+ product study participants.
- **Automated** data-entry **validation** (including retroactive checks) and test scheduling with **Python**, **JavaScript**, and **Google Cloud Platform**, improving data quality and projected to speed downstream **ML** model delivery by 30%.

Projects

Spotify 'Smart Queue' Simulated A/B Experiment | Python, SQL, BigQuery, Mixpanel, Next.js, Databricks Oct 2026

- Designed an end-to-end **A/B experiment** framework for a simulated **feature launch**, pre-registering **OEC, guardrail metrics**, and decision rules. Ran **power analysis** and **SRM checks**, then analyzed a 14,420-user simulation (398K rows) with **Welch**, non-inferiority, novelty-decay, and **bootstrap CACE** methods to deliver a **ship/no-ship recommendation**.
- Defined **tracking schema** (events, properties) for a mockup in Next.js, then built a pipeline from ingestion through **BigQuery SQL** (ported to **Databricks**) to a **self-serve Mixpanel dashboard** covering **funnel**, adoption, and cannibalization metrics.

Amazon GPU ETL Pipeline | Airflow, Docker, Python, PostgreSQL Feb 2025

- Built a scheduled **ETL** (**Airflow** DAG in **Docker**) that scrapes and transforms GPU listings into Postgres for automated ingestion.

Partner Engineering – EDA & Supervised Model | SQL, R, sklearn, pandas Jan 2024 – Jul 2024

- Deployed a **Balanced Random Forest** classifier (winning 3rd place, UCI Data Science Expo) after cleaning and analyzing 40M+ real estate records, contributing to a **Machine Learning** recommendation system to automate a report-writing tool.

Education & Certifications

University of California, Irvine Sep 2024
B.S. Data Science *Donald Bren School of Information and Computer Sciences*

Google Analytics 4 (GA4) Certification *Google Skillshop*